

The ANNIGMA-Wrapper Approach to Fast Feature Selection for Neural Nets

Chun-Nan Hsu, Hung-Ju Huang, and Dietrich Schuschel

Abstract—This paper presents a novel feature selection approach for backprop neural networks (NNs). Previously, a feature selection technique known as the *wrapper model* was shown effective for decision trees induction. However, it is prohibitively expensive when applied to real-world neural net training characterized by large volumes of data and many feature choices. Our approach incorporates a weight analysis-based heuristic called *artificial neural net input gain measurement approximation (ANNIGMA)* to direct the search in the wrapper model and allows effective feature selection feasible for neural net applications. Experimental results on standard datasets show that this approach can efficiently reduce the number of features while maintaining or even improving the accuracy. We also report two successful applications of our approach in the helicopter maintenance applications.

Index Terms—Curse of dimensionality, feature selection, neural networks (NNs), wrapper model.

I. INTRODUCTION

Selection of relevant features is of primary importance to the success of a neural net. The goal is to find the minimum subset of features that yield the highest accuracy. This problem is especially severe when real-world applications are attempted and human selection of features is not available, desirable, or dependable. There are two models of feature subset selection. In the filter model, the features are filtered independently of the induction algorithm. This filtering is done as a pre-processing step. In contrast, the wrapper model [1] wraps around the induction algorithm, searching the feature subset space guided by the performance of the induction algorithm.

Since the filtering model ignores the effect of the feature subset on the performance of the induction algorithm, many researchers have pointed out that it may not be as effective and general as the wrapper model [1]–[3]. They make the point that feature subset selection must take into account the biases of the induction algorithm in order to perform well. However, since in the wrapper model, a large number of training is required to search for the best performing feature subset, it can be prohibitively expensive for neural net applications. Many search strategies were proposed to speed up the search, including hill climbing [3], compound operators [2], randomized algorithms [4], etc. However, when applied to neural net application, at each branching point of the search, these approaches still need to train m neural nets with cross validation to select the next feature subset, where m is the branching factor. This can be prohibitively expensive in real-world applications. A directed feature subset search is needed for neural nets. In this paper's approach, the feature subset search is accelerated by a heuristic called *artificial neural net input gain measurement approximation (ANNIGMA)*. ANNIGMA ranks neural net features by relevance. Following each training, the ANNIGMA heuristic ranks the fea-

tures by relevance. This makes it unnecessary to train m neural nets for each branching point. A huge improvement in speed is now realized. In real-world applications, fewer features means fewer costs to build sensors and to run systems. This paper also reports two successful real-world applications of the ANNIGMA-wrapper approach in helicopter maintenance.

II. THE ANNIGMA-WRAPPER APPROACH

A. An Approximate Metric for Total Gain

The heuristic called ANNIGMA ranks features by relevance based on the weights associated with the features. The reasoning behind this heuristic is that neural net weights can be viewed as representing the gain of the input signal to the output node. Input signals that are noisy or irrelevant to the output will have a high error rate if they have high associated weights. Therefore, training algorithms must reduce their weights such that they do not contribute to the output. In a similar manner, the weights of relevant and noise-free signals will be increased.

The equation for a two-layer neural net with the first layer having a logistic activation function $S(x) = (1/1 + \exp(-x))$ and the second layer having a linear output is

$$O_k = L_k \times \sum_j S\left(\sum_i A_i \times W_{ij}\right) \times W_{jk}$$

where i, j, k are the input, hidden, and output layers node indexes, respectively. L is the second layer linear multiplier value; A is the input node (feature); O is the output node; and W is the weight between the layers. The output O_k as a function of a single input A_i can be expressed as

$$O_k = L_k \times \sum_j S(A_i \times W_{ij} + C_{ij}) \times W_{jk} \quad (1)$$

where C_{ij} represents the constant value of all the other inputs, including biases. C_{ij} here acts as a "setpoint" on the logistic function curve. This equation, with all of the inputs processed through the logistic function, is too complex to analyze directly. An approximation can be made of the total relative gain G of a particular input node i to a particular output node j . The approximation substitutes a linear factor for the logistic activation function. The approximation's error is reduced when the inputs are all in the same range. If we substitute a linear factor F for the activation function, we have

$$O_k \cong L_k \times \sum_j F \times (A_i \times W_{ij} + C_{ij}) \times W_{jk}. \quad (2)$$

The local gain LG is defined to be

$$LG_{ik} = \left| \frac{\Delta O_k}{\Delta A_i} \right|. \quad (3)$$

Since L_k and F are common factors to ANNIGMA's numerator and denominator, they can be dropped in the calculation of LG , i.e.

$$LG_{ik} = \sum_j |W_{ij} \times W_{jk}|. \quad (4)$$

The ANNIGMA score is the local gain LG normalized to a scale of 100

$$\text{ANNIGMA}_{ik} = \frac{LG_{ik}}{\max(LG_k)} \times 100. \quad (5)$$

When the input features are scaled in the same range, the weights will give an approximation of the relative gain of each of these features.

Manuscript received October 3, 2000; revised May 15, 2001, August 28, 2001, and November 6, 2001. This paper was recommended by Associate Editor D. Cook.

C.-N. Hsu is with the Institute of Information Science, Academia Sinica, Nankang 115, Taipei City, Taiwan, R.O.C. (e-mail: chunnan@iis.sinica.edu.tw; http://chunnan.iis.sinica.edu.tw).

H.-J. Huang is with the Department of Computer and Information Science, National Chiao-Tung University, Hsinchu 300, Taiwan, R.O.C.

D. Schuschel is with Longbow Apache Software Boeing Company, Mesa, AZ 85215-9797 USA.

Publisher Item Identifier S 1083-4419(02)00702-1.

TABLE I
SAMPLE RECORDS OF SYNTHESIZED DATA SET

Random	XOR		Sum < 2			Target
A_1	A_2	A_3	A_4	A_5	A_6	O_1
Rand(0,1)	1	1	1	1	1	0
Rand(0,1)	0	0	1	1	0	0
Rand(0,1)	1	0	0	0	1	1
Rand(0,1)	0	1	0	1	0	1
Rand(0,1)	0	1	1	0	0	1




Fig. 1. ANNIGMA output by training iteration.

B. An Illustrative Example of the ANNIGMA Heuristic

We synthesized a small dataset to test and illustrate the predictive power of ANNIGMA heuristic. The dataset for this example has binary valued features that satisfy the normalization requirement of the ANNIGMA heuristic. The dataset contains six columns of input features (A_1 – A_6) and one column of target outputs (O_1). Table I shows sample records of our synthesized data.

Under uncorrupted conditions, the target output O_1 and the meta-signals XOR and SUM < 2 must agree, and other combinations of features are not possible [i.e., $XOR(A_2, A_3) = (A_4 + A_5 + A_6 < 2) = O_1$]. To simulate situations where noise exists along with redundancy, the fourth feature (A_4) was corrupted by flipping the bit with a 20% probability [i.e., the value of A_4 is replaced by its complement with twenty percent probability (i.e., $P[A_4 \leftarrow not(A_4)] = 0.20$)]. This tests the ability of the heuristic to select the noise-free features.

The ANNIGMA scores of these features after six iterations of training epochs is given in Fig. 1. For each of the six features (feature 1–6), there are six bars representing the ANNIGMA scores, G, for each of six neural net training iterations. The rightmost (black) bar of each feature reflects a fully trained net and is the final score. Fig. 1 shows that within a few training iterations, ANNIGMA correctly suppressed the noisy inputs A_1 and A_4 . ANNIGMA correctly identified the XOR signal (A_2 and A_3) as the most relevant. We obtained the same results in 30 of 30 different neural net initializations.

C. Integrating ANNIGMA Into the Wrapper Algorithm

Selecting features solely based on a weight-based metric may yield unreliable results. The ANNIGMA-wrapper approach integrates the ANNIGMA heuristic into the wrapper model to achieve reliable results. Fig. 2 gives an overview of the generic feature selection algorithm in the ANNIGMA-wrapper approach. The input to this algorithm is the set of all features, and their associated training data. The output of this algorithm is a feature subset.

On the top level, this algorithm consists of two nested loops. The outer loop, called *the wrapper cycle*, selects the next subset to evaluate

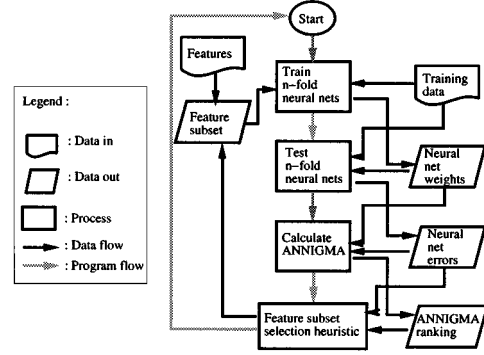


Fig. 2. Program and data flow in the ANNIGMA-wrapper algorithm.

based on the ANNIGMA evaluation of each feature and the classification performance of the neural network (NNs) using this subset of features. How candidate feature subsets are actually generated and evaluated is realized by the box “feature subset selection heuristic” at the bottom of Fig. 2. Feature subset selection heuristic is instantiated by a search strategy which will be described in detail in the next section.

The inner loop called *the training cycle* corresponds to the first and second boxes on the program flow path in Fig. 2. Each of the training cycles takes the candidate feature set as the input and estimates the error rate of a trained neural net using the test data subset. The error rate is estimated by tenfold cross validation as follows. There is an additional “holdout” set of data for final performance testing which is not used here.

The next box on the path of the program flow is to calculate ANNIGMA rankings of the features. After each run of the cross validation is completed, (5) is evaluated for each feature as described in Section II-A. The resulting ANNIGMA score is then weighted by the test error rate [i.e., $ANNIGMA_score * (1 - test_error_rate)$]. This ensures that the better performing neural nets have proportionately greater influence on the final ranking of features. The resulting error-weighted ANNIGMA scores are averaged for n runs to produce the final score for each feature and the ANNIGMA ranking. We compared the results of weighted and nonweighted cases in our experiments. As expected, the weighted case performed better [5].

D. Search Strategies

Our algorithms are based on the strategy of backward elimination [1], [3]. Backward elimination starts with a complete set of original features and removes features from candidate subsets during the search. We present three versions of backward selection including 1) greedy backward elimination (BE), 2) backward elimination with backtracking (BEB), and 3) backward stepwise elimination (BSE) for integrating the ANNIGMA heuristic into the wrapper model.

BE is a greedy version of backward elimination that runs the training cycle to obtain the ANNIGMA rankings and the error rates with all features, then starts with a set including all features. It repeatedly eliminates the next worst ANNIGMA ranked feature in each wrapper cycle until the error rate goes up.

BEB is a version of BE that allows for backtracking. The idea is that if the error rate goes up, instead of terminating the feature selection, the previous feature being eliminated is restored and the next worst ranked feature is eliminated. The process is iterated until a performance-improving elimination is found for each size of feature subsets.

BSE (Algorithm 1) is designed to accelerate feature selection for large datasets when BEB is too slow. The main idea is to eliminate a large number of seemingly irrelevant features in early cycles and adjust

the feature subset carefully in the subsequent cycles. When the performance degrades, the best of the discarded features are brought back into the candidate subset.

Algorithm 1 [Backward Stepwise Elimination (BSE)]

```

01. Run a training cycle with all feature, calculate ANNIGMA ranking  $A(0)$ 
02. Let  $\text{error}(0)$  = averaged error rate and the feature subset  $f(0)$  = all features
03. Let  $H$  = the set of discarded features, initially empty
04. Let cycle counter  $t = 1$ 
05. WHILE NOT termination condition
06.   IF  $t < 4$  THEN LET  $f(t)$  = top  $p_1\%$  of the best features in  $A(t-1)$ 
07.   ELSE IF  $\text{error}(t-1) \leq$  minimum of previous errors THEN
08.     Let  $f(t)$  = top  $p_2\%$  of the best features in  $A(t-1)$ 
09.   ELSE IF  $\text{error}(t-1) \leq$  mean of previous errors THEN
10.     Let  $f(t)$  = top  $p_3\%$  of the best features in  $A(t-1)$  plus the best feature in  $H$ 
11.   ELSE IF  $\text{error}(t-1) \leq$  maximum of previous errors THEN
12.     Let  $f(t)$  = top  $p_4\%$  of the best features in  $A(t-1)$  plus the best 2 features in  $H$ 
13.   ELSE Let  $f(t)$  = top  $p_5\%$  of the best features in  $A(t-1)$  plus the best  $|f(t-1)/2|$  features in  $H$ 
14. Update  $H$ ; Sort  $H$  over the history-averaged ANNIGMA scores
15. Run a training cycle with  $f(t)$ , calculate ANNIGMA ranking  $A(t)$ 
16.  $t = t + 1$ 
17. END WHILE and RETURN  $f(t)$ 

```

The parameters ($p_1 \sim p_5$) used in Algorithm 1 are tuned for the experiment in the pulse-echo classification domain ($p_1 = 75$, $p_2 = 85$, $p_3 = p_4 = 90$, and $p_5 = 50$; see Section V). The reduction percentages (lines 6, 8, 10, 12, and 13) and initial fixed reduction period (line 6) can be adjusted for a given application. Changes to these parameters may affect the efficiency of the search as well as the selected features. In our experiments, we found that slight changes to these parameters did not affect the resulting sets of selected features.

History-averaged score is used as a more reliable estimate of the relevance of a discarded feature than the most recent score because reintroducing a discarded feature is based on an assumption that its relevance is underestimated by the most recent score. However, which score in the previous wrapper cycles is the most reliable is not known, while recomputing the score may incur huge overhead. Therefore, a conservative choice is to use the history-averaged score.

III. EXPERIMENTAL RESULTS AGAINST UCI DATASETS

This section reports experiments exploring the results of using the ANNIGMA-wrapper algorithm against standard artificial and real-world datasets. These datasets were obtained from the UCI Machine-Learning Repository [6].

TABLE II
INFORMATION FOR EACH DATASET

data set	data size	configuration	Algorithm
3P	100+40	8+6+logsig+purelin	BEB
CorrAL	64+64	10+8+tansig+purelin	BE
Monk3a	122+432	15+6+tansig+purelin	BE
Monk3b	122+432	8+3+purelin+purelin	BE
Cancer	399+300	10+12+logsig+logsig	BEB
Credit	490+200	10+10+tansig+purelin	BE
Heart(LB)	133+67	10+2+purelin+purelin	BEB
Ionosphere	200+151	10+22+logsig+purelin	BSE
Pima	576+192	10+6+tansig+logsig	BEB
Vote	218+217	8+3+purelin+purelin	BE

A. Dataset Description and Preparation

The first column of Table II gives the size of our experimental datasets (in the format: training/cross validation + holdout set.) The ratio of training and holdout set is 2:1 unless explicitly designated by dataset providers. Some datasets contain missing values. We simply replace all missing values by the feature mean for all instances belonging to the same class. The first four rows are artificial datasets and the rest are real-world datasets.

B. Experimental Procedure

The second column of Table II describes the neural net configurations (in the format: maximum epochs + number of hidden nodes + hidden layer transfer function + output layer transfer function) for each dataset in the experiments. These configurations are empirically determined based on their cross validation errors and relative variances of the resulting ANNIGMA rankings.

The experiments are carried out as follows. For each dataset, the first step is to normalize input features into the same range, usually between 0 and 1. After the configuration is determined, we estimate the error rate of the neural net with no feature selection by applying tenfold cross validation to train ten sets of the weights using the training set, and then testing them against the holdout set and averaging the resulting error rates. Next, we compare the feature selection performance of ANNIGMA-wrapper by applying different search strategies described in Section II-D 30 times and report the average number of features selected, the average error rate, and the average execution time. In each trial, the elements in training set and holdout set of a testing dataset are randomly selected except the dataset whose training and holdout set have been explicitly designated by its provider. The error rate is estimated by averaging holdout set errors after the final feature subset is selected. Two different search strategies are applied to each dataset for performance comparisons: 1) SWB and 2) BE family. SWB is the standard wrapper-backward elimination. It is a greedy version of backward elimination for the standard wrapper feature selection, where error rates are the sole metric to guide the search. This approach is introduced here for the purpose of comparisons. The fourth column of Table II shows which BE family algorithm is used for each dataset. We use BSE for datasets with many features, BE for small or synthesized datasets, and BEB for others, mainly to maximize the utility of our computing facilities. The experiments are conducted on Pentium III 800 MHz PC's with 128 MB memory.

TABLE III
PERFORMANCE STATISTICS OF FEATURE SELECTION BY ANNIGMA-WRAPPER

data set	NN		SWB			BE family			RD	Recent advances		
	F.	Error%	F.	Error%	Time	F.	Error%	Time		F.	Error%	Ref.
3P	13	9.3±14.7	4.4±3.1	0.5±1.4	1165	3.0±0.0	0.0±0.0	175	5.3	3.0±0.0	0.0±0.0	hybrid
CorrAL	6	2.3±4.0	4.0±0.2	0.0±0.2	98	4.0±0.2	0.0±0.2	50	1.4	4.0±0.0	12.5±0.0	INFO
Monk3a	6	10.0±5.2	3.4±1.6	5.1±3.4	125	2.3±0.7	2.9±0.8	58	2.2	—	—	—
Monk3b	15	2.8±0.0	4.4±1.1	2.8±0.0	519	2.2±0.4	2.8±0.0	56	0.0	3.9±1.8	1.6±1.7	NNFS
Cancer	9	4.1±4.7	7.2±1.2	3.6±1.1	451	5.8±1.3	3.5±1.2	280	1.8	2.7±1.0	5.9±1.0	NNFS
Credit	15	14.1±1.7	13.4±1.0	14.4±0.8	1712	6.7±2.5	12.0±0.8	375	0.7	7.0±0.0	14.9±6.1	AHOC
Heart(LB)	13	20.2±0.2	7.3±1.9	25.7±1.8	352	2.7±1.2	22.3±2.0	121	1.0	5.0±0.0	19.2±6.5	AHOC
Ionosphere	34	11.4±3.9	32	10.2	138065	9.0±2.5	9.8±1.3	342	1.3	—	—	—
Pima	8	24.1±5.0	6.9±1.0	23.0±1.3	179	5.2±1.4	22.2±1.4	168	1.9	2.9±0.2	25.7±3.3	NNFS
Vote	16	3.2±0.0	3.2±1.4	3.2±0.2	1052	3.3±1.9	3.1±0.2	65	0.1	5.0±0.0	4.3±3.5	AHOC

C. The Results

Table III reports the average execution time in seconds to complete a feature selection task for each dataset with different algorithms. In all cases, the ANNIGMA-wrapper algorithms (BE family) can complete a feature selection task in less than 10 min. compared with SWB, the ANNIGMA-wrapper algorithms are many times faster for all datasets, especially for those datasets with a large number of features, such as Ionosphere (34 features).

Table III also reports the feature selection performance of the ANNIGMA-wrapper approach. The “NN” column lists the results with no feature subset selection. The second column (“SWB”) reports the results of the standard wrapper-backward elimination. The third column compare the results by using the search strategies (BE family) of the ANNIGMA-wrapper approach. For each case (except for SWB for Ionosphere), we report the average and the standard deviation of the number of features selected, the average and the standard deviation of the error rates. Since it takes too much time for SWB to select features for dataset Ionosphere, we only complete this algorithm one time for Ionosphere and report only that result.

We also use ANOVA with the Bonferroni procedure for multiple comparisons statistics [7]. The difference between any two error rates in a row must be at least as large as the value in the “Required Difference” or “RD” column in order to be considered statistically significant at the 90% confidence level for the experiment as a whole. An error rate in boldface is significantly better than that of “NN.”

D. Discussion

The results show that BE family perform well in general for seven out of the ten datasets, improving the number of features or the error rate over the base NN significantly. The results for 3P and CorrAL are particularly remarkable. BEB selects the correct three features for 3P and achieves perfect classification in all 30 tests. Even the replicated features are distinguished and eliminated. For CorrAL, BE almost always selects the correct four features, occasionally adding an irrelevant one.

The results also show that feature expansion can have a negative effect when it does not offset the cost of adding another feature. For Monk3a, using features as-is, the correct features were selected. For Monk3b, the correct features were chosen, but never the minimum subset. Note that when using the expanded representation, there are aliases for each feature (e.g., $a_4 = 1$ implies $a_4 \neq 2$ and $a_4 \neq 3$).

Comparing the results for Monk3a with Monk3b, we see that ANNIGMA-wrapper can identify relevant features even when they are multiple-valued. We also tried to perform feature expansion for some datasets (such as Credit) which have nonbinary categorical features, but results, not reported here, were time-consuming and the performances are not improved significantly.

IV. COMPARISONS WITH RECENT RELATED WORK

Recently, many clever techniques have been proposed to improve the effectiveness of feature selection for neural nets, ranging from filter model-based approaches to genetic algorithms. This section compares their results with ours.

A. Recent Advances

Richeldi and Lanzi presented an approach called AHOC [8], which partitions the observed features into a number of groups, called factors, that reflect the major dimensions of the phenomenon under consideration. A genetic algorithm is used to explore the feature space originated by the factors and to determine the set of the most informative feature configurations.

Neural network feature selector (NNFS) [9] is a method that adds a penalty term to the error function used to derive the weight updating rule of NN training.

GADistAl [10] is a wrapper-based approach to feature selection using a genetic algorithm in conjunction with a constructive NN learning algorithm called DistAl [11], which is employed to evaluate the “fitness” of candidate feature subsets in the genetic algorithm. The “fitness” function is designed to combine the classification accuracy and the cost of using a set of features.

Dash and Liu [12] proposed a hybrid algorithm of probabilistic search [13] and complete search to take advantage of both algorithms. It begins with Las Vegas filter (LVF) [4], a probabilistic feature selection algorithm, to reduce the number of features and then runs Automatic Branch & Bound (ABB), a complete search algorithm.

B. Comparisons

The algorithms that are surveyed for the comparison including the original wrapper [1], Las Vegas filter (LVF) [13], Las Vegas wrapper (LVW) [4], neural net feature selector (NNFS) [9], hybrid approach

(hybrid) [12], information-theoretic filter (INFO) [14], and AHOC genetic algorithm (AHOC) [8].¹ The last column of Table III shows these performance data. The comparison is made on the basis of both the number of features selected and the error rates after feature selection. It should be mentioned that results achieved by different algorithms are not obtained with the same experimental procedure, nor did we implement and rerun their algorithms but using their report literally. Therefore, the comparison is inherently informal.

According to the performance data, among eight datasets, ANNIGMA-wrapper is better in six cases (3P, CorrAL, Cancer, Credit, Pima, and Vote) in terms of the error rates. Hybrid [12] also achieves perfect result for 3P, but their 3P dataset contains 12 features while we use a 3P dataset with 13 features. For datasets Monk3b, and Heart (LB), ANNIGMA-wrapper produces the smallest feature subset but higher error rates than the best of recent advances. However, we note that the training dataset of Monk3 has 5% class noise, and the holdout dataset contains the training dataset. The 12 noisy records out of a total of 432 records result in an expected best error rate of 2.78%. This is exactly the error rate achieved by ANNIGMA-wrapper in most of its cases. An induction algorithm achieving a better error rate might *overfit* the data by modeling the noise.

V. APPLICATIONS IN THE HELICOPTER INDUSTRY

This section reports two successful applications of the ANNIGMA-wrapper approach in the helicopter industry.

A. Pulse-Echo Experiment

This system classifies ultrasonic pulses that are used to inspect parts formed of laminated composite materials. In this system, the inspection process consists of sending a short ultrasonic pulse into the material under the test. These echo signals are collected and digitized by an A/D converter. Any change in the materials propagation speed for sound within the sample will result in an echo being sent back toward the transducer that initiates the test pulse. Material failures are associated with echo signal shapes.

The digitized waveform is converted into 192 features. The system must classify these digitized waveforms into “good” and “bad” categories. An existing system uses a wavelet-compression approach to select top 25 features.

In this experiment, the ANNIGMA-wrapper algorithm uses all 192 features as its initial feature set. One cycle takes about 1 min on a Pentium-Pro 200MHz equipped computer. Fig. 3(a). shows the features selected as the ANNIGMA-wrapper algorithm progressed for material “45b45.” Fig. 3(b). shows the number of features selected by a wrapper cycle. The associated performance is presented in Fig. 3(c). The accuracies are based on 150 pulse-echo waveform that were not used in the training. When this ANNIGMA-wrapper algorithm was run, stopping criteria was turned off. It can be seen that cycles after cycle 48 are not necessary.

The existing system, using the top 25 wavelet coefficients, achieved 97% accuracy. The ANNIGMA-wrapper algorithm, used only ten of these coefficients to achieve 99% accuracy. The intersection of the two sets of coefficients contains only two members. The result suggests that the set of wavelet coefficients contains a large number of redundant relevant features. The ANNIGMA-wrapper approach manages to select a smaller set of features while achieving a slightly better accuracy.

¹Table III does not include the results of GADistAI because they reported tenfold cross validation results rather than the results for holdout data. Their results may not be comparable with other approaches.




Fig. 3. ANNIGMA-wrapper results by cycle for material 45b45.

B. Strain Signal Prediction

The second application is the development of a system that predicts the strain on helicopter rotor blades. Helicopter rotor blade repair and replacement are major cost factors in helicopter operations. However, they are essential to assure helicopter safety. Current industry standard to measure the fatigue life is based on a fixed number of operating hours assuming most severe operation. This is an inefficient because perfect blades may be mistakenly scrapped. A proposed solution is to track the strains on the blades to calculate fatigue damage [15], [16]. However, it is not feasible to directly measure strains on production helicopters due to the high expense of a sensor/monitoring system, and the unreliability of a strain sensor mounted on a rotor blade. We attempt to solve this problem through a neural net trained to predict strain from other easily obtainable helicopter signals. The neural nets in this application is to approximate the strain gauge reading mounted on the helicopter blades as a function of the given sensory inputs. Initially, 41 sensors from among thousands available on the helicopter were selected by flight dynamics experts. We have 197 sets of training datasets (1G bytes large total). Each set contains the sensory data collected from 3- to 19-s periods.

Due to very large size of the datasets, the first wrapper cycle takes 200 h of exclusive run-time on a 200 MHz Pentium Pro-based computer. After three wrapper cycles, ANNIGMA-Wrapper selects 19 features from 41 and reduce the error rate from 17.3% to 16.3%. The fourth cycle yields 13 features but the error rate jumps to 77.4%. The implication of these results is that the information content in the 13 selected features is insufficient compared to that in the 19 features. These 19 signals are therefore recommended for further system development. The first wrapper cycle was rerun to check the reliability of the ANNIGMA heuristic. We found that the top ten high-ranking features were within two ranks of differences between two runs [5].

VI. CONCLUSIONS

In this paper, we have presented a new approach to selecting features for neural nets called ANNIGMA-wrapper that makes the wrapper model feature selection tractable in real-world neural net applications. Experimental results against standard datasets from UCI repository show that our simple approach performs well for datasets with various characteristics, and the ANNIGMA-wrapper approach was successfully applied to two real-world neural net applications in the helicopter industry.

ACKNOWLEDGMENT

The authors wish to thank Y.-T. Yang for her help in programming our experiments and anonymous reviewers for their valuable comments.

REFERENCES

- [1] G. H. John, R. Kohavi, and K. Pfleger, "Irrelevant features and the subset selection problem," in *Machine Learning: Proceedings of the Eleventh International Conference (ICML '94)*. San Mateo, CA: Morgan Kaufmann, 1994.
- [2] R. Kohavi and D. Sommerfield, "Feature subset selection using the wrapper method: Overfitting and dynamic search space topology," in *Proceedings of the First International Conference on Knowledge Discovery and Data Mining (KDD '95)*. Menlo Park, CA: AAAI Press, 1995.
- [3] R. Carunan and D. Freitag, "Greedy attribute selection," in *Machine Learning: Proceedings of the Eleventh International Conference*. San Mateo, CA: Morgan Kaufmann, 1994.
- [4] H. Liu and R. Setiono, "Feature selection and classification—A probabilistic wrapper approach," in *Proc. 9th Int. Conf. Industrial and Engineering Applications of AI and ES*, 1996.
- [5] D. Schuschel, "Automation of attribute selection for neural nets," M. S. Thesis, Dept. Computer Science and Engineering, Arizona State Univ., Tempe, 1998.
- [6] C. J. Merz and P. M. Murphy. UCI repository of machine learning databases. Dept. Computer Science, Univ. California, Irvine. [Online]. Available: <http://www.ics.uci.edu/~mllearn/MLRepository.html>.
- [7] J. T. McClave and F. H. Dietrich, *Statistics*. San Francisco, CA: Dellen, 1991.
- [8] M. Richeldi and P. L. Lanzi, "Performing effective feature selection by investigating the deep structure of the data," in *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD '96)*, E. Simoudis, J. Han, and U. M. Fayyad, Eds. Menlo Park, CA: AAAI Press, 1996.
- [9] R. Setiono and H. Liu, "Neural network feature selector," *IEEE Trans. Neural Networks*, vol. 8, pp. 654–662, May 1997.
- [10] J. Yang and V. Honavar, "Feature subset selection using a genetic algorithm," in *Feature Extraction, Construction, and Subset Selection: A Data Mining Perspective*, H. Motoda and H. Liu, Eds. Norwell, MA: Kluwer, 1998, ch. 8.
- [11] J. Yang, R. Parekh, and V. Honavar, "DistAL: An inter-pattern distance-based constructive learning algorithm," *Intell. Data Anal.*, vol. 3, pp. 55–73, 1999.
- [12] M. Dash and H. Liu, "Hybrid search of feature subsets," in *Proc. PRICAI*, Singapore, Nov. 1998.
- [13] H. Liu and R. Setiono, "A probabilistic approach to feature selection—A filter solution," in *Proceedings of the 13th International Conference on Machine Learning (ICML '96)*. Bari, Italy, 1996, pp. 319–327.
- [14] D. Koller and M. Sahami, "Toward optimal feature selection," in *Machine Learning: Proc. of the Thirteenth International Conference (ICML '96)*, L. Saitta, Ed. San Mateo, CA: Morgan Kaufmann, 1996.
- [15] S. Downing and D. Socie, "Simple rainfall counting algorithms," *Int. J. Fatigue*, Jan. 1982.
- [16] *Standard Practices for Cycle Counting in Fatigue Analysis*, American Society for Testing and Materials, ASTM E ASTM 1049-85, 1985.

Granular Clustering: A Granular Signature of Data

Witold Pedrycz and Andrzej Bargiela

Abstract—The study is devoted to a granular analysis of data. We develop a new clustering algorithm that organizes findings about data in the form of a collection of information granules—hyperboxes. The clustering carried out here is an example of a granulation mechanism. We discuss a compatibility measure guiding a construction (growth) of the clusters and explain a rationale behind their development. The clustering promotes a data mining way of problem solving by emphasizing the transparency of the results (hyperboxes). We discuss a number of indexes describing hyperboxes and expressing relationships between such information granules. It is also shown how the resulting family of the information granules is a concise descriptor of the structure of the data—a granular signature of the data. We examine the properties of features (variables) occurring of the problem as they manifest in the setting of the information granules. Numerical experiments are carried out based on two-dimensional (2-D) synthetic data as well as multivariable Boston data available on the WWW.

Index Terms—Complex systems, confidence limits analysis, data mining, feature analysis, granular time series, hyperboxes, information abstraction, information granules and granulation, interval analysis, principle of balanced information granularity.

I. INTRODUCTORY COMMENTS

Making sense of data has been a motto of data mining. Any in-depth analysis of data that leads to comprehensive and interpretable results has to address an issue of transparency of final findings. In one way or another, arises a need for casting the results in the language of information granules—conceptual entities that capture the essence of the overall data set in a compact manner. It is worth stressing that information granules are a vehicle of abstraction that supports a conversion of clouds of numeric data into more tangible information granules [2], [3], [5], [12], [13], [16]–[18].

The area of clustering with its long history has been an important endeavor of finding structures in data and representing the essence of such finding in terms of prototypes, dendrograms, self-organizing maps [8], [9] and alike [1], [4]. Commonly, if not exclusively, the direct aspect of granulation has not been tackled. The intent of this study is to address this important problem by introducing an idea of granular clustering. Being more descriptive, the simplest scenario looks like this: we start from collection of numeric data (points in \mathbf{R}^n) and form information granules whose distribution and size reflects the essence of the data. Forming the clusters (information granules) may be treated as a process of growing information granules—as the clustering progresses, we expand the clusters, enhance the descriptive facet of the granules while gradually reduce the amount of details being available to us. The information granules we are interested in this study are represented as hyperboxes positioned in a highly dimensional data space. The mathematical formalism of the interval analysis provides a robust framework for the analysis of information density of the granular structures that

Manuscript received February 3, 2001; revised November 6, 2001. This work was supported by the Natural Sciences and Engineering Research Council of Canada (NSERC), Alberta Consortium of Software Engineering (ASERC) and by the Engineering and Physical Sciences Research Council (U.K.). This paper was recommended by Associate Editor P. K. Bhattacharya.

W. Pedrycz is with the Department of Electrical and Computer Engineering, University of Alberta, Edmonton, AB, Canada and also with the Systems Research Institute, Polish Academy of Sciences 01-447 Warsaw, Poland (e-mail: pedrycz@ee.ualberta.ca).

A. Bargiela is with the Department of Computing, The Nottingham Trent University, Nottingham NG1 4BU, U.K. (e-mail: andre@doc.ntu.ac.uk).

Publisher Item Identifier S 1083-4419(02)00703-3.